

SAI PAVAN ADITYA MANTRIPRAGADA

📞 +918179367134 📩 saipavanadityam@gmail.com

LinkedIn: SAI PAVAN ADITYA MANTRIPRAGADA GitHub: SAI-PAVAN-ADITYA-MANTRIPRAGADA

Summary

Strategic AI/ML Engineer with a proven track record of architecting scalable data ecosystems and high-impact analytics solutions that drive enterprise-level decision-making. Expert in bridging the gap between robust data engineering (Snowflake, SnapLogic) and advanced ML research (Federated Learning, Differential Privacy). Proactive leader focused on transforming legacy descriptive workflows into predictive engines, directly enhancing operational efficiency and stakeholder ROI. Committed to delivering privacy-first, production-grade AI while driving cross-functional collaboration to leverage data as a core strategic asset.

Education

University of Essex <i>Master of Science in Data Science (Distinction)</i>	April 2024 Colchester, UK
<ul style="list-style-type: none">Specialization: Advanced Machine Learning, Large-Scale Data Engineering, and Algorithmic AI Ethics.Research Thesis: Developed a Privacy-Preserving ML framework using Differential Privacy & Federated Learning to ensure GDPR compliance in decentralized datasets.Achievement: Graduated with Distinction; Specialized in cloud-native AI deployment strategies.	
Osmania University <i>Bachelor of Engineering in Computer Science (Distinction)</i>	July 2021 Hyderabad, India
<ul style="list-style-type: none">Core Competencies: Data Structures & Algorithms, Distributed Systems, and Artificial Intelligence.Capstone Project: Augmented Reality Indoor Navigation—Integrated Computer Vision (OpenCV) with ARKit for real-time spatial mapping (Published Research).	

Technical Skills

Languages & Core: Python (Expert), SQL (Advanced), R, Bash, C++
AI & Generative Models: LLMs (GPT-4, Llama 3), Fine-tuning (PEFT/LoRA), RAG Pipelines, LangChain, Transformers, Computer Vision (OpenCV)
Machine Learning & Deep Learning: PyTorch, TensorFlow, Scikit-Learn, XGBoost, Differential Privacy, Federated Learning
MLOps & Production: CI/CD for ML, Model Monitoring (Drift/Bias), MLflow, DVC, Weights & Biases, Flask/FastAPI
Data Engineering & Cloud: Snowflake, SnapLogic, Spark, AWS (SageMaker, Lambda, S3), Docker, Kubernetes, Terraform (IaC)
Data Governance & Analytics: Vector Databases (Pinecone, Milvus), Differential Privacy Frameworks, Power BI, Advanced Statistical Modeling

Experience

Zensar Technologies <i>AI – ML Engineer</i>	June 2025 – Present Hyderabad, India
<ul style="list-style-type: none">Strategic Data Architecture: Architected end-to-end automated data pipelines using SnapLogic and Python to ingest high-velocity Salesforce CRM data into Snowflake, reducing data latency and accelerating executive reporting cycles by 40%.Enterprise BI Strategy: Spearheaded the development of a real estate analytics suite, delivering high-fidelity dashboards that provide cross-functional stakeholders with actionable insights into portfolio performance and market trends.Predictive Roadmap Leadership: Leading the technical transition from descriptive metrics to predictive modeling frameworks, deploying ML-driven forecasting to optimize asset management and drive long-term strategic investment outcomes.System Integration: Engineered a seamless data bridge between Salesforce and the proprietary 'Inspace' application, centralizing disparate data streams into a unified source of truth to empower real-time, data-driven operational decisions.	
Greene King <i>Inventory & Supply Chain Data Analyst</i>	April 2024 – May 2025 London, UK
<ul style="list-style-type: none">Demand Forecasting: Deployed an ML-driven forecasting engine using AWS SageMaker and Time-Series analysis, resulting in a 7% reduction in overstock costs and significantly improving supply chain resilience across regional hubs.	

- **Revenue Optimization:** Developed sophisticated customer segmentation models utilizing **Scikit-Learn**, which refined targeted marketing strategies and directly contributed to an **18% increase in campaign conversion rates**.
- **Churn Mitigation:** Architected a predictive retention framework using **XGBoost and Random Forest** to identify at-risk accounts, successfully reducing customer churn by 12% through proactive, data-informed intervention strategies.
- **Process Automation:** Optimized internal data workflows by replacing manual SQL/Python scripts with automated ETL triggers, slashing report generation time from **120 minutes to under 10 minutes** for the supply chain leadership team.

Cognizant

Programmer Analyst

March 2021 – July 2022

Chennai, India

- **System Reliability:** Managed critical enterprise content management systems (Documentum, xPlore), maintaining a **95% resolution rate** for complex technical issues within strict SLA parameters for global clients.
- **Workflow Engineering:** Redesigned document rendition workflows to improve conversion speeds, achieving a **30% reduction in processing errors** through systematic optimization of DocShifter and Insight platforms.
- **SAP Integration:** Facilitated the synchronization of SAP-Documentum metadata via the ASSAP framework, automating manual data entry processes and increasing overall data integrity by reducing human intervention by 40%.

Projects

Enterprise RAG Engine for Legal & Patent Intelligence | *LangChain, Pinecone, Llama-3, FastAPI, Docker*

- Architected a **Retrieval-Augmented Generation (RAG)** system to automate the summarization of 500k+ patent documents, utilizing **Vector Databases** for sub-second semantic search and context retrieval.
- Fine-tuned Large Language Models using **QLoRA** to improve legal terminology accuracy by 40%, significantly reducing the "hallucination" rate compared to out-of-the-box foundation models.
- Deployed the solution as a microservice-based architecture, reducing manual document review time by 60% and enabling scalable, real-time query handling via an optimized API for legal researchers.

Privacy-Preserving AI Framework (Production-Grade) | *PySyft, Federated Learning, Differential Privacy*

- Developed a decentralized training framework to enable model training on sensitive financial and healthcare datasets while maintaining strict **GDPR and SOC2 compliance** through Differential Privacy mechanisms.
- Implemented **Federated Averaging (FedAvg)** algorithms, allowing multiple data silos to contribute to a global model without raw data exchange, maintaining a 92% accuracy rate compared to centralized training.
- Packaged the solution into a reusable Python library for data anonymization, providing automated noise-injection and epsilon-budget tracking to prevent membership inference attacks in production environments.

MLOps Pipeline for Real-Time Predictive Maintenance | *MLflow, AWS SageMaker, Kubernetes, DVC*

- Engineered a robust **MLOps lifecycle** to automate the deployment of predictive models, integrating **CI/CD pipelines** that handle automated retraining upon detecting data drift or performance degradation.
- Utilized **Kubernetes** for container orchestration and AWS SageMaker for scalable inference, ensuring the system could handle a 300% surge in request volume while maintaining low-latency response times.
- Reduced technical debt by implementing **Data Version Control (DVC)**, ensuring 100% reproducibility of experiments and enabling seamless rollbacks of model versions in complex enterprise environments.

Autonomous Inventory Optimization Engine | *PyTorch, Transformer-based Time Series, AWS Lambda*

- Designed a high-dimensional forecasting engine using **Temporal Fusion Transformers**, outperforming traditional ARIMA/LSTM models by 25% in capturing seasonal demand fluctuations for supply chain logistics.
- Integrated a serverless inference layer using **AWS Lambda**, optimizing cloud compute costs by 30% while delivering real-time stock-out risk scores to regional warehouse management dashboards via Power BI.
- Directly influenced operational efficiency by reducing overstocking costs by 15%, providing business stakeholders with a probabilistic view of inventory needs through an interactive and scalable interface.

Leadership & Recognition

Operations Head: Established the **Matrusri Developer Space**, a technical community of 200+ members; orchestrated workshops on Cloud Computing and AI to bridge the industry-academia gap.

Open Source Contributor: Actively developing a Python-based utility for **Differential Privacy** to assist in GDPR-compliant data masking (Upcoming Publication tie-in).

Awarded: Nominated for **Best Student Award** by the Institution of Engineers (India) for excellence in technical leadership and research innovation (2021).

Certifications

Microsoft Certified: Azure AI Engineer Associate (AI-102): Validated expertise in building, managing, and deploying AI solutions using Azure Cognitive Services and Azure OpenAI.

Microsoft Power BI Data Analyst (PL-300): Certified in designing enterprise-grade BI solutions, advanced data modeling, and delivering actionable insights for stakeholders.

Microsoft Certified: Azure Data Scientist Associate (DP-100): Specialized in designing and implementing data science solutions and machine learning workloads on the Azure platform.

Snowflake Pro Core Certification: Advanced proficiency in Snowflake cloud architecture, performance tuning, and scalable data warehouse optimization in multi-cloud environments.

Publications & Research

Federated Learning for Decentralized Data Privacy | In Press (2026): Developed a novel framework for cross-silo model training that minimizes communication overhead while maintaining 90%+ accuracy.

Adaptive Differential Privacy in Financial Datasets | In Press (2026): Researching epsilon-budget optimization to balance data utility and privacy in high-dimensional financial records.

Augmented Reality Indoor Navigation System | IRJMETS, Vol. 3 (2021): Published research on computer-vision-based localization using OpenCV and feature extraction for offline navigation.